

Lab 2

Kevin Liang : kg1392

Matthew Tan : mmt2338

Nicholas Sehy : nws326

1, $Z \sim N(0, 1)$, $n = 10,000$ comp. $P(Z_{avg} > 0.01)$

$$Z = \frac{\bar{X} - \mu}{\sigma / \sqrt{n}} \quad \leftarrow \text{population mean}$$

 $\sigma = \text{pop. std. dev.}, N = \text{samples} = 10,000$

$$\frac{0.1 - 0}{1/\sqrt{10,000}} = \frac{0.1}{1/100} = 10 \quad P = 0\%$$

$$\frac{0.01 - 0}{1/100} = 1 \quad P(Z < 1) = 0.8413$$

$$\frac{0.01 - 0}{1/100} = 0.1 \quad P(Z < 0.1) = 0.5398$$

$$\begin{aligned} & \frac{1}{\sqrt{n}} \sim \frac{1}{\sqrt{10,000}} \\ & \frac{1}{\sqrt{n}} \sim \frac{1}{100} \end{aligned}$$

$$CLT: \frac{Z - \mu}{\sigma / \sqrt{n}} \quad \frac{(n^{-1/2})}{\sigma / \sqrt{n}} = \frac{1}{\sigma}$$

$$\frac{n^{-1/2}}{\sigma / \sqrt{n}} = \frac{n^{1/2}}{\sigma}$$

$$\frac{n^{-1/2}}{\sigma / \sqrt{n}} = \frac{n^{1/2}}{\sigma}$$

$$P(Z_{avg} - \mu > n^{-1/m}) = P\left(Z < \frac{n^{1/m} - 1/2}{\sigma}\right)$$

$$Z, y_i = x_i^T \beta + e_i \quad \min_{\beta} : \frac{1}{n} \sum_{i=1}^n (x_i^T \beta - y_i)^2$$

$$\sum_{i=1}^n (x_i^T \beta - y_i)^2 = \frac{1}{n} \sum_{i=1}^n x_i^T \beta^2 - 2 x_i^T y_i \beta + y_i^2 = 0$$

$$A = x_i^T \quad B = -2 y_i x_i \quad C = y_i^2$$

A is positive (x_i is some scalar so x_i^2 is always > 0)

$$A \beta^2 + B \beta + C = 0 \Rightarrow \beta^T x^T x \beta - 2 \beta^T x^T y + y^T y$$

$$\nabla (\beta^T x^T x \beta) = 2 x^T x \beta$$

$$\nabla \beta^T x^T y = x^T y$$

$$\nabla y^T y = 0$$

$$A^T A x = A^T b \quad \beta = A^{-1} b$$

$$\nabla (\|x\beta - y\|_2^2) = 2 x^T x \beta - 2 x^T y + 0 = 0$$

$$\hat{\beta} = (x^T x)^{-1} x^T y$$

$$y_i = x_i^T \beta + e_i, \quad \hat{\beta} = \beta + z e$$

e = noise/error vector
 z = matrix

$$\hat{\beta} = [(x^T x)^{-1} x^T] [y_i^T \beta + e_i^T]$$

$$\hat{\beta} = (x^T x)^{-1} x^T (x_i^T \beta) + (x^T x)^{-1} x^T e^T$$

$$z = (x^T x)^{-1} x$$

```
In [113]: import numpy as np
import pandas as pd
import sklearn as sk
import matplotlib.pyplot as plt
import seaborn as sns
import math
```

```
In [114]: # Question 1 - Part 1
a = pd.read_csv("D:\Kevin Liang\Downloads\Lab2_Data\DF1", header = 0, index_co
l = 0)

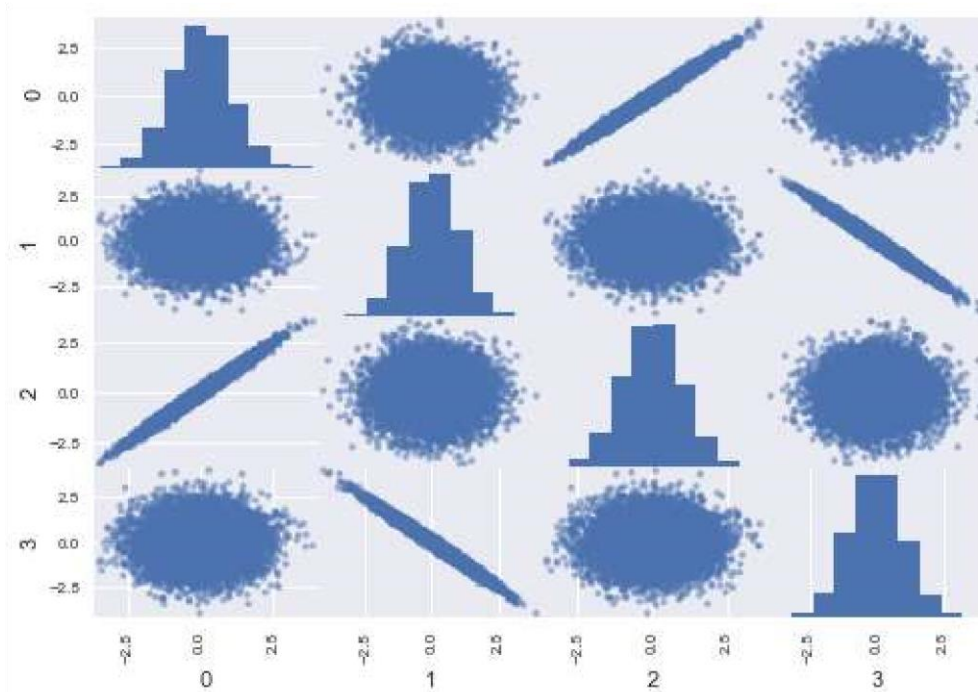
#pandas
print a.corr()

b = pd.plotting.scatter_matrix(a)

#seaborn
sns.pairplot(a,kind='scatter')
sns.plt.show()


# Column's 0 and 2 are pairwise correlated with a correlation value of 0.99006
6.
# Column's 1 and 3 are pairwise correlated with a negative correlation value o
f -.990235.
```

	0	1	2	3
0	1.000000	-0.003998	0.990066	0.004111
1	-0.003998	1.000000	-0.004085	-0.990235
2	0.990066	-0.004085	1.000000	0.004067
3	0.004111	-0.990235	0.004067	1.000000



	0	1	2	3
--	---	---	---	---

0	1.001558	-0.004012	0.991624	0.004125
1	-0.004012	1.005378	-0.004099	-0.995457
2	0.991624	-0.004099	1.001589	0.004081
3	0.004125	-0.995457	0.004081	1.005168



In [115]: *# Question 1 - Part 2*

`a.cov()`

Covariance Equation : $(1/n) \cdot \text{SUM}((x - \text{mean}(x))(y - \text{mean}(y)))$

The positive elements of the covariance matrix are also positive

in the correlation matrix and in the plots, while the negative elements



are also negative in the correlation matrix and in the correlation plots.

Out[115]:


```
In [116]: # Question 1 - Part 3

cov = []
for x in range(2,1000):
    c = np.random.multivariate_normal([0,0,0], [[1,0,0],[0,1,.8],[0,.8,1]], x)
    d = np.cov(c.transpose())
    cov.append(d[1,2])

plt.scatter(range(len(cov)), cov)
plt.show()
plt.plot(range(len(cov)), cov)
plt.show()
```



```
In [117]: #Problem 2

data = pd.read_csv("D:\Kevin Liang\Downloads\Lab2_Data\DF2")


# displaying current data showing two outliers
plt.scatter(data['0'], data['1'])
plt.show()

# determing covariance matrix of data set
a = np.cov(data['0'], data['1'])


# taking inverse of data
a = np.linalg.inv(a)
xyData = data.loc[:, '0':'1']

# multiplying inverse covariance matrix times every data point to get a gaussian distribution of data,
# and to get the covariance matrix of the data to be the identity matrix
k = np.dot(xyData, a)
print k

#displaying data with one huge outlier point (1,1) transformed
plt.scatter(k[:,0], k[:,1])
plt.show()
```



```
[ [ 10.4672271  -9.53225769]
  [  8.54640488 -8.31353788]
  [-10.55682902 10.72311449]
  ...,
  [-8.47775263  9.46233257]
  [-2.78440519  3.29159129]
  [-0.55553807  1.77574979]]
```



```

In [118]: # Problem 3

# Part 1

# Declare constants
def part_two(n):
    beta_naught = -3
    beta = 0
    samples = 2000

    distances = []
    for x in range(samples):
        x = np.random.normal(0, 1, n)
        e = np.random.normal(0, 1, n)

        # Generate  $y_i = \beta_{naught} + e_i$ 
        y = beta_naught + x * beta + e

        # Calculate  $\beta_{hat}$  for the dataset
        beta_hat = np.dot(x, y) / np.dot(x, x)
        #  $\beta_{hat} = (x.transpose() * y) * (x.transpose() * x)$ 
        sd = beta_hat - beta
        distances.append(sd)

    sum = 0
    for num in distances:
        sum += math.pow(num, 2)
    std_dev = math.sqrt(sum / samples)
    return std_dev

print part_two(150)

# For  $n = 150$ , the empirical standard deviation is 0.263219887789.
# This means that the Beta Hat which we find (-0.15) is not at all significant
# as this is less than the deviation.

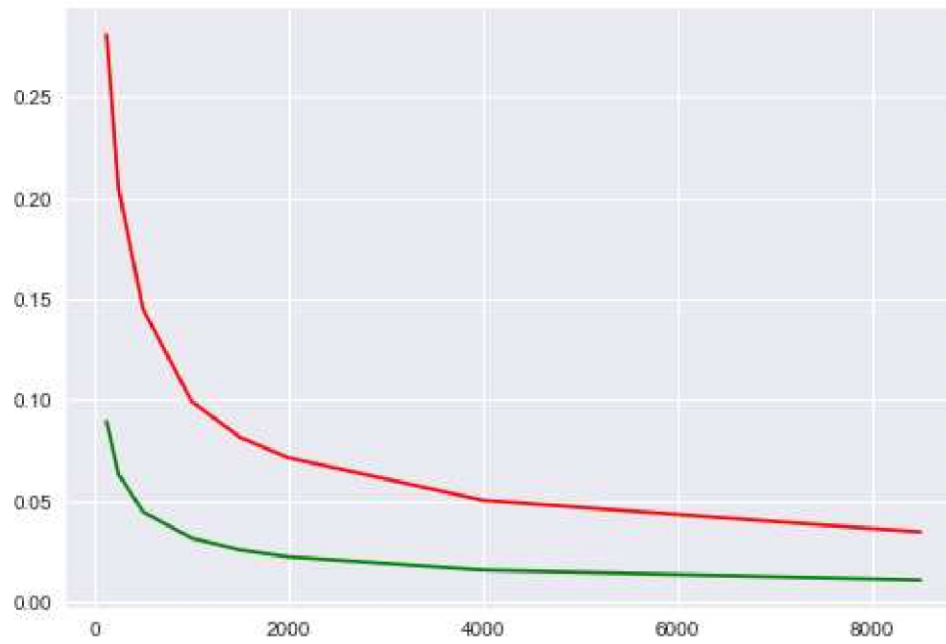
# Part 2
n = [125, 250, 500, 1000, 1500, 2000, 4000, 8500]
plot_n = []
for number in n:
    plot_n.append(part_two(number))
plot_sqrt = []
for number in n:
    temp = np.sqrt(number)
    temp = 1/temp
    plot_sqrt.append(temp)
print(plot_sqrt)
plt.plot(n, plot_n, 'r')
# plt.plot(plot_n)
plt.plot(n, plot_sqrt, 'g')
plt.show()

# the fit for both of the graphs are similar in form. They both decrease at an
# exponential rate.

```

0.263219887789

[0.0894427190999991588, 0.063245553203367583, 0.044721359549995794, 0.031622776601683791, 0.025819888974716113, 0.022360679774997897, 0.015811388300841896, 0.010846522890932807]



```

In [119]: # Problem 4

# Part 1
def most_popular(k, XXXX):
    df = pd.read_csv('D:\Kevin Liang\Downloads\Names\Names\yob' + str(XXXX) +
".txt", header=None, names=["Name", "Gender", "Frq"], delimiter=',', usecols=
[0,2])
    df = df.groupby(by=["Name"]).sum()
    #names.sort_values(3, ascending=False)
    print df.nlargest(int(k), "Frq")
    #print names

# Part 2
def name_frequency(Name):
    men_frq = 0
    women_frq = 0
    num_adds = 0
    for year in range(1880, 2016):
        df = pd.read_csv('D:\Kevin Liang\Downloads\Names\Names\yob' +
str(year) + ".txt", header=None, names=["Name", "Gender", "Frq"],
delimiter=',')
        select_indices = list(np.where(df["Name"] == str(Name))[0])
        if Name in df["Name"].values:
            for _ in select_indices:
                if df["Gender"].iloc[_] == 'M':
                    men_frq += df["Frq"].iloc[_]
                    num_adds += 1
                else:
                    women_frq += df["Frq"].iloc[_]
                    num_adds += 1
        else:

```

0

```

        print Name + " wasn't found in year " + str(year)
    print str(men_frq) + " instances of men named " + str(Name)
    print str(women_frq) + " instances of women named " + str(Name)

# Part 3
def relative_frequency(Name, year):
    men_frq = 0.0
    women_frq = 0.0
    num_adds = 0
    #for year in range(1880, 2016):
    df = pd.read_csv('D:\Kevin Liang\Downloads\Names\Names\yob' + str(year) +
".txt", header=None, names=["Name", "Gender", "Frq"], delimiter=',')
    select_indices = list(np.where(df["Name"] == Name)[0])
    total_freq = 0.0
    for val in df["Frq"]:
        total_freq += int(val)
    if Name in df["Name"].values:
        for _ in select_indices:
            if df["Gender"].iloc[_] == 'M':
                men_frq += df["Frq"].iloc[_]/total_freq
            else:
                women_frq += df["Frq"].iloc[_]/total_freq
    else:
        print Name + " wasn't found in year " + str(year)
    print "Normalized frequency of " + Name + " in men:" + str(men_frq)
    print "Normalized frequency of " + Name + " in women:" + str(women_frq)

# Part 4
def gendered_popularity():
    D = { }
    a = [ ]
    men_frq = { } #largest value only
    women_frq = { } #largest value only
    num_adds = 0
    winner = ' '
    for year in range(1880, 2016):
        df = pd.read_csv('D:\Kevin Liang\Downloads\Names\Names\yob' +
str(year) + ".txt", header=None, names=["Name", "Gender", "Frq"],
delimiter=',')
        size = len(df)
        name = ""
        for index in range (0, size):
            name = str(df["Name"][index])
            if not(men_frq.has_key(name)):
                men_frq[name] = 0
            if not(women_frq.has_key(name)):
                women_frq[name] = 0;
            if df["Gender"][index] == 'M' and men_frq[name] < df["Frq"]
[index]:
                men_frq[name] = df["Frq"][index]
            if D.has_key(name) and D[name] != 'M':
                if women_frq.has_key(name) and women_frq[name] < men_frq[n
ame] and women_frq > 0:
                    D[name] = 'M'
                    a.append(name)
            else:
                D[name] = 'M'

```

```

        else:
            D[name] = ' '
    else:
        women_frq[name] = df["Frq"][index]
        if D.has_key(name) and D[name] != 'F':
            if women_frq.has_key(name) and women_frq[name] > men_frq[n
ame] and men_frq > 0:
                D[name] = 'F'
                a.append(name)
            else:
                D[name] = 'F'
        else:
            D[name] = ' '
    my_set = set(a)
    print len(my_set)

# TEST
most_popular(10,1880)
name_frequency("Matt")
relative_frequency("John", 1880)
gendered_popularity()

```

```

    Frq
Name
John    9701
William 9561
Mary    7092
James   5949
Charles 5359
George  5152
Frank   3255
Joseph  2642
Anna    2616
Thomas  2542
23460 instances of men named Matt
5 instances of women named Matt
Normalized frequency of John in men:0.0479199134414
Normalized frequency of John in women:0.000228308235971
69479

```



```
In [120]: # Problem 5

# Method 1 - Displayed
# My method consists of checking all user_location cells and determine whether
# or not
# the name of the state exists somewhere in the user_locations cell.
# This is done by checking if the substring (the state name) exists in the user
#_location cell.
# It will then mark a tweet with its associated state.
# The total number of tweets from a state is then calculated and displayed in
# a list.

# Method 2 - Not Displayed
# Regular expressions could be used to account for state acronyms as well.

tweets = pd.read_csv("D:/Kevin Liang/Downloads/tweets.csv")
```

3

```
# dictionary of all states in the US
states = {
    'AK': 'Alaska',
    'AL': 'Alabama',
    'AR': 'Arkansas',
    'AS': 'American Samoa',
    'AZ': 'Arizona',
    'CA': 'California',
    'CO': 'Colorado',
    'CT': 'Connecticut',
    'DC': 'District of Columbia',
    'DE': 'Delaware',
    'FL': 'Florida',
    'GA': 'Georgia',
    'GU': 'Guam',
    'HI': 'Hawaii',
    'IA': 'Iowa',
    'ID': 'Idaho',
    'IL': 'Illinois',
    'IN': 'Indiana',
    'KS': 'Kansas',
    'KY': 'Kentucky',
    'LA': 'Louisiana',
    'MA': 'Massachusetts',
    'MD': 'Maryland',
    'ME': 'Maine',
    'MI': 'Michigan',
    'MN': 'Minnesota',
    'MO': 'Missouri',
    'MP': 'Northern Mariana Islands',
    'MS': 'Mississippi',
    'MT': 'Montana',
    'NA': 'National',
    'NC': 'North Carolina',
    'ND': 'North Dakota',
    'NE': 'Nebraska',
    'NH': 'New Hampshire',
    'NJ': 'New Jersey',
    'NM': 'New Mexico',
    'NV': 'Nevada',
    'NY': 'New York',
    'OH': 'Ohio',
    'OK': 'Oklahoma',
    'OR': 'Oregon',
    'PA': 'Pennsylvania',
    'PR': 'Puerto Rico',
    'RI': 'Rhode Island',
    'SC': 'South Carolina',
    'SD': 'South Dakota',
    'TN': 'Tennessee',
    'TX': 'Texas',
    'UT': 'Utah',
    'VA': 'Virginia',
    'VT': 'Vermont',
    'WA': 'Washington',
    'WI': 'Wisconsin',
    'WV': 'West Virginia',
    'WY': 'Wyoming'
}
```

```

        'WI': 'Wisconsin',
        'WV': 'West Virginia',
        'WY': 'Wyoming'
    }

    # area equation to determine check which state of a tweet
    def area(location):
        for x in states.keys():
            if states[x].lower() in str(location).lower():
                return states[x]

    # create new column categorizing each tweet with a associated state
    tweets["place"] = tweets["user_location"].apply(area)

    # counting the total number of tweets from a state
    t1 = {}
    t1 = tweets["place"].value_counts()

    # display count of tweets from a state
    print t1

```

The print result isn't showing.

Texas	4795
New York	4520
Florida	4042
California	4007
Washington	3335
Ohio	1326
New Jersey	1207
North Carolina	1138
Michigan	1138
Colorado	1070
Georgia	1031
Virginia	1008
Tennessee	990
Arizona	986
Indiana	910
Pennsylvania	890
Oregon	822
Kansas	734
Illinois	713
Maryland	579
Alabama	572
South Carolina	559

Massachusetts	521
Wisconsin	514
Minnesota	509
Oklahoma	504
Louisiana	486
Missouri	462
Iowa	447
Kentucky	397
Connecticut	367
Nevada	312
Maine	258
West Virginia	249
Hawaii	247
Vermont	242
New Hampshire	238
Utah	211
Delaware	199
Montana	176
Alaska	173
National	172
Idaho	158
Mississippi	153
New Mexico	151
Nebraska	145
Rhode Island	120
Puerto Rico	113
Wyoming	55
Virgin Islands	36
North Dakota	26
District of Columbia	25
South Dakota	24
Guam	7
American Samoa	1

Name: place, dtype: int64